

# Review of Inverse Reinforcement Learning frameworks for Imitation Learning

Josselin Bonnevie  
Shahine Bouabid

MVA

## 1 Introduction

- What is Imitation Learning ?
  - Behavioral Cloning
  - Direct policy learning
- Problem Setup

## 2 A Linear Programming Approach

## 3 Max-Margin algorithms

- A Quadratic Programming approach
- Max-margin Planning


## 4 Maximum Entropy Inverse Reinforcement Learning

- MaxEnt Principle
- Deep IRL
- Generative Adversarial Imitation Learning

# Introduction




Bojarski et al. [2]



Taylor et al. [5]

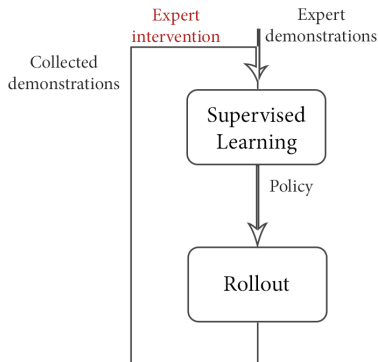
**Given** demonstrations  $\rightarrow$  **Find** followed policy

# Introduction



- Poor behavior recovery if step deviation
- Quadratic loss in task horizon

# Introduction



- Linear loss in task horizon
- Requires interactive feed-back

*Idea :* Use reward learning as a mean for policy derivation

# Problem Setup

## Framework :

- MDP :  $\mathcal{M} = (S, A, P, R^*, \gamma, I)$  endowed with expert policy  $\pi^*$
- Energy :  $\mathcal{E}(\pi, R) = \mathbb{E}_{s \sim I}[V^\pi(s)]$

## Given :

- $\mathcal{M} \setminus R^*$
- Environment simulation
- Demonstrations  $\mathcal{D} = \{\zeta_i^*\}_i$  where  $\zeta_i^* \sim \pi^*$
- Energy evaluation

## Objective :

$$\text{Find } R \text{ s.t. } \pi^* \in \arg \max_{\pi} \mathcal{E}(\pi, R)$$

**Approximation :** Leverage  $\mathcal{D} \rightarrow \text{Find } R \text{ s.t. } \pi_{\mathcal{D}} \in \arg \max_{\pi} \mathcal{E}(\pi, R)$

# Problem Setup

**Environment feature functions :**  $\phi = (\phi_j)$ ,  $\phi_j : S \times A \rightarrow \mathbb{R}$

- Model environment facets
- Introduce prior knowledge


→ *Linear reward parametrization* :  $R = \theta^\top \phi$

**Expected feature count :**

$$\mu(\pi) = \mathbb{E} \left[ \sum_{t \geq 0} \gamma^t \phi(s_t, a_t) \mid \pi \right] \Rightarrow \mathcal{E}(\pi, \theta) = \theta^\top \mu(\pi)$$

# A Linear Programming Approach

**Idea :** Favor reward functions making deviation wrt optimal policy costly



Given  $\{\pi_1, \dots, \pi_k\}$  and linear penalization  $h$ , maximize :

$$\sum_{i=1}^k h(\mathcal{E}(\pi_{\mathcal{D}}, \theta) - \mathcal{E}(\pi_i, \theta)), \quad |\theta_j| \leq 1 \forall j \quad (\text{LP}_k)$$



# A Linear Programming Approach

---

**Algorithm 1** Ng and Russell [3]


---

- 1: Initialize  $k \leftarrow 1$  ;  $\theta$  randomly ;  $\mathcal{E}(\pi_{\mathcal{D}}, \theta) \leftarrow \theta^{\top} \mu(\mathcal{D})$
  - 2: Initialize base case policy  $\pi_1$  and collect trajectories  $\mathcal{Z}_1$
  - 3: Initialize  $\mathcal{E}(\pi_1, \theta) \leftarrow \theta^{\top} \mu(\mathcal{Z}_1)$
  - 4: **while** Termination condition not met **do**
  - 5:    $\theta \leftarrow \text{solve (LP}_k\text{)}$
  - 6:   Derive  $\pi_{k+1} = \arg \max_{\pi} \mathcal{E}(\pi, \theta)$
  - 7:   Roll out to collect trajectories  $\mathcal{Z}_{k+1}$
  - 8:   Update  $\mathcal{E}(\pi_{\mathcal{D}}, \theta) \leftarrow \theta^{\top} \mu(\mathcal{D})$  and  $\mathcal{E}(\pi_j, \theta) \leftarrow \theta^{\top} \mu(\mathcal{Z}_j) \forall j \in \{1, \dots, k\}$
  - 9:    $k \leftarrow k + 1$
  - 10: **end while**
  - 11: **return**  $\pi_k$
- 

→ Simple, intuitive but no convergence guarantee

# A Quadratic Programming approach

**Idea :** Move problem to feature expectation space :  $\mathcal{E}(\pi, \theta) = \theta^\top \mu(\pi)$   
with constraint  $\|\theta\|_2 \leq 1$



$$\max_{\eta, \theta} \quad \eta$$

$$\text{s.t.} \quad \theta^\top \mu(\mathcal{D}) - \eta \geq \theta^\top \mu(\mathcal{Z}_i) \quad \forall i \in \llbracket 1, k \rrbracket \text{ and } \|\theta\|_2 \leq 1$$

(SVM<sub>k</sub>)

# A Quadratic Programming approach

---

**Algorithm 2** Abbeel and Ng [1]

---

- 1: Initialize  $\eta \leftarrow 0$  ;  $k \leftarrow 1$  ;  $\theta$  randomly
  - 2: Initialize base case policy  $\pi_1$  and collect trajectories  $\mathcal{Z}_1$
  - 3: Compute  $\mu(\mathcal{D})$  and  $\mu(\mathcal{Z}_1)$
  - 4: **while**  $\eta > \varepsilon$  **do**
  - 5:    $\eta, \theta \leftarrow \text{solve}(\text{SVM}_k)$
  - 6:   Derive  $\pi_{k+1} = \arg \max_{\pi} \mathcal{E}(\pi, \theta)$  under  $R$
  - 7:   Roll out to collect trajectories  $\mathcal{Z}_{k+1}$  and compute  $\mu(\mathcal{Z}_{k+1})$
  - 8:    $k \leftarrow k + 1$
  - 9: **end while**
  - 10: **return**  $\pi_k$
-

# A Quadratic Programming approach

## Theorem

Let  $\varepsilon > 0$ ,  $\delta \in ]0, 1[$ ,  $d$  the dimensionality of  $\theta$  and  $m$  the number of Monte Carlo samples used.

- If  $m \geq \frac{2d}{(1-\gamma)^2\varepsilon^2} \log \frac{2d}{\delta}$  and  $\|\theta\|_2 \leq 1$
- Then  $\exists K = \mathcal{O}\left(\frac{d}{(1-\gamma)^2\varepsilon^2} \log \frac{d}{(1-\gamma)\varepsilon}\right)$  s.t.  $|\mathcal{E}(\pi^*, \theta) - \mathcal{E}(\pi_K, \theta)| \leq \varepsilon$  with probability  $1 - \delta$

## Pros :

- Convergence in a finite number of steps
- Control over solution's accuracy

## Cons :

- Strong assumption : existence of a near-optimal policy
- Lack of generalization for start/goal variations

# Max-margin Planning

## Idea :

- Map each demonstration  $\zeta_i^*$  to a different MDP endowed with its own expert policy  $\pi_i^*$  and set of feature functions
- But unique weight vector  $\theta$
- Introduce for each MDP an occupancy loss vector  $\ell_i \rightarrow$  weights in closeness to  $\pi_i^*$

$$\begin{aligned} \min_{\eta, \theta} \quad & \frac{\lambda}{2} \|\theta\|_2^2 + \frac{\kappa}{n} \sum_{i=1}^n \beta_i \eta_i^q \\ \text{s.t.} \quad & \theta^\top \mu(\zeta_i^*) + \eta_i \geq \max_{\zeta} \theta^\top \mu(\zeta) + \ell_i^\top \rho^\zeta \quad \forall i \end{aligned} \quad (\text{MMP}_k)$$

Number of constraints  $\propto |\mathcal{D}| \prod_i |S_i \times A_i| \rightarrow$  subgradient resolution

---

**Algorithm 3** Ratliff et al. [4]

---

```
1: Initialize hyperparameters, learning rate  $\alpha_t$ , horizon  $T$ ,  $t \leftarrow 1$ ,  $\theta \leftarrow 0$ 
2: while  $t \leq T$  do
3:   for  $i$  in  $1 \dots n$  do
4:     Update  $R_i \leftarrow \theta^\top \phi(i)$ 
5:     Derive  $\hat{\pi}_i^* = \arg \max_{\pi} \mathcal{E}(\pi, R_i) + \mathcal{L}_i(\pi)$ 
6:   end for
7:   Update  $\theta \leftarrow \theta - \alpha_t \cdot \text{subgradient}(\hat{\pi}_1^*, \dots, \hat{\pi}_n^*)$  and increment  $t$ 
8: end while
9: return  $\theta$ 
```

---

**Great :**

- Weight vector prone to extrapolate to new feature maps
- Convergence guaranteed in batch and online settings (resp. linear and sublinear)

# Feature and occupancy matching

## Expected feature count

$$\mu(\pi) = \mathbb{E} \left[ \sum_{t \geq 0} \gamma^t \phi(s_t, a_t) \mid \pi \right] \Rightarrow \mathcal{E}(\pi, \theta) = \theta^\top \mu(\pi)$$

## Occupancy measure

$$\rho^\pi(s, a) = \mathbb{E}_\pi \left[ \sum_{t \geq 0} \gamma^t \mathbf{1}_{(s_t, a_t) = (s, a)} \right] \Rightarrow \mathcal{E}(\pi, \theta) = \sum_{(s, a)} \rho^\pi(s, a) R(s, a) = R^\top \rho^\pi$$

*Idea :* Find policy with same performance for whatever the reward

# Maximum Entropy Principle - Deterministic MDP

Many policies can match the the expected features of the expert.

## Entropy Maximization on the trajectory distribution

$$\max_{p_\pi} H(p_\pi) \quad \text{s.t.} \quad \mu(p_\pi) = \hat{\mu}(\mathcal{D})$$

## Maximum likelihood in the exponential family

$$\hat{p}(\zeta) = \frac{e^{\theta^\top \mu(\zeta)}}{Z(\theta)} \quad \text{with} \quad Z(\theta) = \int e^{\theta^\top \mu(\zeta)} d\zeta$$

## Gradient

$$\nabla_\theta L(\theta) = \hat{\mu}(\mathcal{D}) - \mu(\pi_\theta) = \hat{\mu}(\mathcal{D}) - \sum_{(s,a) \in S \times A} \rho^\theta(s,a) \phi(s,a)$$



# Maximum Causal Entropy for non Deterministic MDPs:

**Causal Entropy ( $\gamma$  discounted):**

$$\tilde{H}^\gamma(\pi) = \mathbb{E} \left[ - \sum_{t \geq 0} \gamma^t \log \pi(a_t | s_t) \right]$$

**The associated IRL problem**

$$\max_{\pi} \tilde{H}^\gamma(\pi) \quad \text{s.t.} \quad \mu(\pi) = \hat{\mu}(\mathcal{D})$$

**Soft Bellman Equations:**

$$\hat{\pi}(a | s) = \exp \left( Q_{\hat{\theta}}^{\text{soft}}(s, a) - V_{\hat{\theta}}^{\text{soft}}(s) \right) \quad (1)$$

**Where**

$$Q_{\hat{\theta}}^{\text{soft}}(s, a) = \theta^\top \phi(s, a) + \beta \sum_{s' \in \mathcal{S}} P(s' | s, a) V_{\hat{\theta}}^{\text{soft}}(s') \quad (2)$$
$$V_{\hat{\theta}}^{\text{soft}}(s) = \operatorname{softmax}_{a \in A} Q_{\hat{\theta}}^{\text{soft}}(s, a) = \log \left( \sum_{a \in A} \exp(Q_{\hat{\theta}}^{\text{soft}}(s, a)) \right)$$

# An other vision of MaxEnt IRL :

Maximum Entropy and Maximum Likelihood are dual problems :

$$\begin{aligned}\hat{\pi} &= \operatorname{argmax}_{\pi} \hat{H}^{\gamma}(\pi) + \hat{\theta}^{\top} (\mu(\pi) - \hat{\mu}(\mathcal{D})) \\ &= \operatorname{argmax}_{\pi} \hat{H}^{\gamma}(\pi) + \mathcal{E}(\pi, \hat{\theta})\end{aligned}$$

Maximum Entropy IRL can somehow be seen as a paradigm with a model on the agent's behaviour.

- We assume that for some  $R^*$  :

$$\pi^* \in \operatorname{argmax}_{\pi} \hat{H}^{\gamma}(\pi) + \mathcal{E}(\pi, R^*)$$

- We want to find  $\hat{R}$  such that

$$\pi^* \in \operatorname{argmax}_{\pi} \hat{H}^{\gamma}(\pi) + \mathcal{E}(\pi, \hat{R})$$

# Maximum Causal Entropy - Guarantees

Trajectories with equal reward have equal probabilities

**Guarantees :**

$$\hat{\pi} \in \arg \min_{\pi} \sup_{\tilde{\pi}} \mathbb{E}_{\tilde{\pi}} \left[ \sum_{t \geq 0} -\beta^t \log \pi(a_t | s_t) \right]$$

## Gradient of the demonstration's likelihood:

$$\nabla_{\theta} L(\theta) = \hat{\mu}(\mathcal{D}) - \mu(\pi_{\theta}) = \hat{\mu}(\mathcal{D}) - \sum_{(s,a) \in S \times A} \rho^{\theta}(s, a) \phi(s, a)$$

## Solving the MDP

- Soft Value Iteration

$$T_{\theta}^{\text{soft}}(V)(s) = \text{softmax}_a \left( \theta^{\top} f(s, a) + \sum_{s' \in S} P(s' | s, a) V(s') \right)$$

- Soft Q-learning :

$$Q_{\theta}^{\text{soft}}(s_t, a_t) \leftarrow Q_{\theta}^{\text{soft}}(s_t, a_t) + \eta(t) \left[ \theta^{\top} f(s_t, a_t) + \gamma \text{softmax}_{a_{t+1}} Q_{\theta}^{\text{soft}}(s_{t+1}, a_t) - Q_{\theta}^{\text{soft}}(s_t, a_t) \right]$$

## Estimating occupancy measures

- Dynamic programming or Monte-Carlo methods

# Deep Inverse Reinforcement Learning

- Feature engineering : hard and painful
- Occupancy measure matching : computing a reward for each action-state pair

Middle ground :

$$R(s, a) = R_{\theta}(\mathbf{f}(s_t, a_t))$$

- Back-propagation of the likelihood's gradient to the network's parameter
- A way to include expert's knowledge

# Generative Adversarial Imitation Learning

- Directly Recovering a policy
- No RL in a loop
- Using Neural Networks

## Causal Entropy Inverse Reinforcement Learning Problem

$$\min_{R \in \mathcal{C}} \left( \max_{\pi \in \Pi} \hat{H}^{\gamma}(\pi) + \mathbb{E}_{\pi}[R(s, a)] \right) - \mathbb{E}_{\pi^*}[R(s, a)]$$

**But we don't want exact occupancy measure matching**

$$\min_{R \in \mathcal{C}} \left( \max_{\pi \in \Pi} \hat{H}^{\gamma}(\pi) + \mathbb{E}_{\pi}[R(s, a)] \right) - \mathbb{E}_{\pi^*}[R(s, a)] \quad \underbrace{+ \psi(R)}_{\text{Regularization}}$$

# Generative Adversarial Imitation Learning

$$RL \circ IRL_{\psi}(\pi^*) = \arg \max_{\pi} \hat{H}^{\gamma}(\pi) - \psi^*(\rho^{\pi} - \rho^{\pi^*})$$

**By taking:**

$$\psi_{GA}(R) = \begin{cases} \mathbb{E}_{\pi^*}[g(R(s, a))] & \text{if } R > 0 \\ +\infty & \text{otherwise} \end{cases}$$

With  $g(R) = R - \log(1 - e^{-R})$  for all  $R > 0$

**We get**

$$\begin{aligned} \psi_{GA}^*(\rho^{\pi} - \rho^{\pi^*}) &= \max_{D \in (0,1)^{S \times A}} \mathbb{E}_{\pi}[\log D(s, a)] + \mathbb{E}_{\pi^*}[\log(1 - D(s, a))] \\ &= D_{JS}(\pi, \pi^*) + cst \end{aligned}$$

# Generative Adversarial Imitation Learning

Final formulation of Generative Adversarial Imitation Learning

$$\text{Find } \hat{\pi} \in \arg \min_{\pi} D_{\text{JS}}(\rho^{\pi}, \rho^{\pi^*}) - \lambda \tilde{H}^{\gamma}(\pi)$$

Or equivalently,

$$\hat{\pi} \in \arg \min_{\pi} \max_{D \in (0,1)^{\mathcal{S} \times \mathcal{A}}} \mathbb{E}_{\pi}[\log(D(s, a))] + \mathbb{E}_{\pi^*}[\log(1 - D(s, a))] - \lambda \tilde{H}^{\gamma}(\pi)$$



- [1] P. Abbeel and A. Y. Ng. Apprenticeship learning via inverse reinforcement learning. In *Proceedings of the Twenty-first International Conference on Machine Learning*, ICML '04, pages 1–, New York, NY, USA, 2004. ACM. ISBN 1-58113-838-5. doi: 10.1145/1015330.1015430. URL <http://doi.acm.org/10.1145/1015330.1015430>.
- [2] M. Bojarski, D. D. Testa, D. Dworakowski, B. Firner, B. Flepp, P. Goyal, L. D. Jackel, M. Monfort, U. Muller, J. Zhang, X. Zhang, J. Zhao, and K. Zieba. End to end learning for self-driving cars. *CoRR*, abs/1604.07316, 2016. URL <http://arxiv.org/abs/1604.07316>.
- [3] A. Y. Ng and S. J. Russell. Algorithms for inverse reinforcement learning. In *Proceedings of the Seventeenth International Conference on Machine Learning*, ICML '00, pages 663–670, San Francisco, CA, USA, 2000. Morgan Kaufmann Publishers Inc. ISBN 1-55860-707-2. URL <http://dl.acm.org/citation.cfm?id=645529.657801>.
- [4] N. D. Ratliff, J. A. Bagnell, and M. A. Zinkevich. Maximum margin planning. In *Proceedings of the 23rd International Conference on Machine Learning*, ICML '06, pages 729–736, New York, NY, USA, 2006. ACM. ISBN 1-59593-383-2. doi: 10.1145/1143844.1143936. URL <http://doi.acm.org/10.1145/1143844.1143936>.
- [5] S. Taylor, T. Kim, Y. Yue, M. Mahler, J. Krahe, A. Garcia Rodriguez, J. Hodgins, and I. Matthews. A deep learning approach for generalized speech animation. *ACM Transactions on Graphics*, 36:1–11, 07 2017. doi: 10.1145/3072959.3073699.